

# ML EXP 12 - YASH ASHOK SHIRSATH TE AIDS-69

```
import pandas as yash
import numpy as shirsath
import seaborn as train
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error

link =
"https://raw.githubusercontent.com/amankharwal/Website-data/master/CarPrice.csv"
carkadata = yash.read_csv(link)
```

## Exploratory Data Analysis

```
carkadata.head()
```

	car_ID	symboling	CarName	fueltype	aspiration
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std

	carbody	drivewheel	engine	location	wheelbase	...
0	convertible	rwd	front	88.6	...	130
1	convertible	rwd	front	88.6	...	130
2	hatchback	rwd	front	94.5	...	152
3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm
citympg \						
0	mpfi	3.47	2.68	9.0	111	5000
21						
1	mpfi	3.47	2.68	9.0	111	5000
21						
2	mpfi	2.68	3.47	9.0	154	5000
19						
3	mpfi	3.19	3.40	10.0	102	5500
24						
4	mpfi	3.19	3.40	8.0	115	5500
18						

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

carkadata.tail()

	car_ID	symboling	CarName	fueltype	aspiration	doornumber
\						
200	201	-1	volvo 145e (sw)	gas	std	four
201	202	-1	volvo 144ea	gas	turbo	four
202	203	-1	volvo 244dl	gas	std	four
203	204	-1	volvo 246	diesel	turbo	four
204	205	-1	volvo 264gl	gas	turbo	four

	carbody	drivewheel	enginelocation	wheelbase	...	enginesize
fuelsystem \						
200	sedan	rwd	front	109.1	...	141
mpfi						
201	sedan	rwd	front	109.1	...	141
mpfi						
202	sedan	rwd	front	109.1	...	173
mpfi						
203	sedan	rwd	front	109.1	...	145
idi						
204	sedan	rwd	front	109.1	...	141
mpfi						

boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
-----------	--------	------------------	------------	---------	---------	---

200	3.78	3.15	9.5	114	5400	23
201	3.78	3.15	8.7	160	5300	19
202	3.58	2.87	8.8	134	5500	18
203	3.01	3.40	23.0	106	4800	26
204	3.78	3.15	9.5	114	5400	19

	highwaympg	price
200	28	16845.0
201	25	19045.0
202	23	21485.0
203	27	22470.0
204	25	22625.0

[5 rows x 26 columns]

carkadata.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 205 entries, 0 to 204

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

dtypes: float64(8), int64(8), object(10)

memory usage: 41.8+ KB

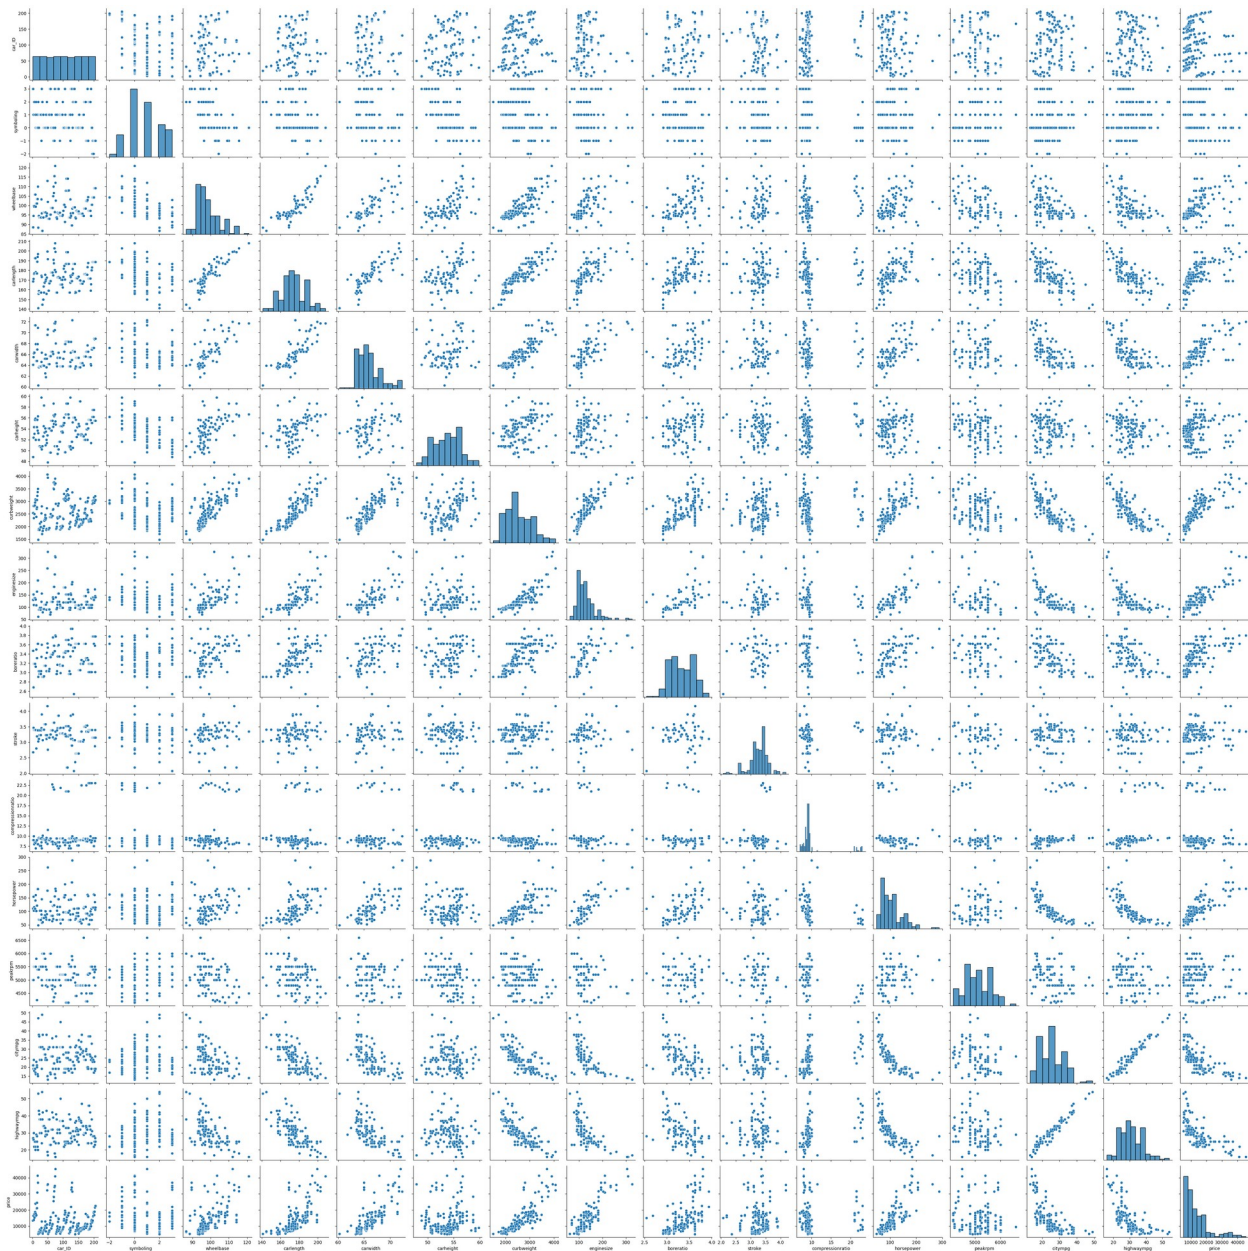
```
carkadata.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth
carheight \					
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805
std	59.322565	1.245307	6.021776	12.337289	2.145204
min	1.000000	-2.000000	86.600000	141.100000	60.300000
25%	52.000000	0.000000	94.500000	166.300000	64.100000
50%	103.000000	1.000000	97.000000	173.200000	65.500000
75%	154.000000	2.000000	102.400000	183.100000	66.900000
max	205.000000	3.000000	120.900000	208.100000	72.300000

	curbweight	enginesize	boreratio	stroke
compressionratio \				
count	205.000000	205.000000	205.000000	205.000000
mean	2555.565854	126.907317	3.329756	3.255415
std	520.680204	41.642693	0.270844	0.313597
min	1488.000000	61.000000	2.540000	2.070000
25%	2145.000000	97.000000	3.150000	3.110000
50%	2414.000000	120.000000	3.310000	3.290000
75%	2935.000000	141.000000	3.580000	3.410000
max	4066.000000	326.000000	3.940000	4.170000

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

```
train.pairplot(carkadata)
plt.show()
```



## Linear Regrsson Modelling

```
# Identify & Remove Non Numeric Data
binanumberwalekolam =
carkadata.select_dtypes(include=['object']).columns
numberwaladata = carkadata.drop(columns=binanumberwalekolam)
```

```

Xash = numberwaladata.drop(columns=['price']) # Xash is Predictor
Yash = numberwaladata['price'] # Yash is Target

Xash_train, Xash_test, Yash_train, Yash_test = train_test_split(Xash,
Yash, test_size=0.2, random_state=42)

Yash_Cha_Model = LinearRegression()
Yash_Cha_Model.fit(Xash_train, Yash_train)

LinearRegression()

Yash_Cha_Prediction = Yash_Cha_Model.predict(Xash_test) # TeSP

# Mean Squared Error measures how close a regression line is to a set
of data points.
mse = mean_squared_error(Yash_test, Yash_Cha_Prediction)
print("Mean Squared Error:-", mse)

Mean Squared Error:- 11710105.078807332

# one of the two main performance indicators for a regression model.
rmse = shirsath.sqrt(mse)
print("Root Mean Squared Error:-", rmse)

Root Mean Squared Error:- 3422.0030798944836

# MAE is a common metric used in statistics and machine learning to
assess the performance of regression models
mae = mean_absolute_error(Yash_test, Yash_Cha_Prediction)
print("Mean Absolute Error:-", mae)

Mean Absolute Error:- 2411.093962409612

# how well the model explains the variability in the dependent
variable based on the independent variables.
r_squared = r2_score(Yash_test, Yash_Cha_Prediction)
print("R-Squared:-", r_squared)

R-Squared:- 0.8516657126363221

```

In the context of linear regression or other regression models, accuracy is not typically used as a performance metric. Instead, we use metrics such as Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and R-squared ( $R^2$ ) to evaluate the model's performance.